

Kev-4B → GPT-OSS-20B cascade: accuracy / latency / cost vs. threshold
filled = simulated from standalone replays (outlined = Pareto-optimal);
open ring = live run; dotted line = threshold in use

